

# Assignment 2

## Autoencoder and Transformer Architectures

### Part I: Theoretical Part

#### Part I.I: Autoencoders for Anomaly Detection

**Task 1: Calculate the total number of trainable parameters in the autoencoder, including weights and biases. Show your detailed steps for each layer.**

The autoencoder consists of multiple fully connected (FC) layers. Each layer has weights and biases.

**Input Layer  $\rightarrow$  Hidden Layer 1 (512 neurons):**

$$\text{Weights} = 1000 \times 512$$

$$\text{Biases} = 512$$

$$\text{Total parameters} = 1000 \times 512 + 512 = 512512$$

**Hidden Layer 1  $\rightarrow$  Bottleneck Layer (32 neurons):**

$$\text{Weights} = 512 \times 32$$

$$\text{Biases} = 32$$

$$\text{Total parameters} = 512 \times 32 + 32 = 16416$$

**Bottleneck Layer  $\rightarrow$  Hidden Layer 2 (512 neurons):**

$$\text{Weights} = 32 \times 512$$

$$\text{Biases} = 512$$

$$\text{Total parameters} = 32 \times 512 + 512 = 16896$$

**Hidden Layer 2  $\rightarrow$  Output Layer (1000 neurons):**

$$\text{Weights} = 512 \times 1000$$

$$\text{Biases} = 1000$$

$$\text{Total parameters} = 512 \times 1000 + 1000 = 513000$$

**Total trainable parameters:**

$$512512 + 16416 + 16896 + 513000 = 1058824$$

**Task 2: Explain how the L2 regularization term will be incorporated into the training process with the MSE loss. Describe its impact on the weight updates during backpropagation.**

The L2 regularization term is added to the loss function to prevent overfitting by penalizing large weight values.

**MSE Loss Function:**

$$\mathcal{L}_{MSE} = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$$

where  $N$  is the number of samples,  $\hat{y}_i$  is the predicted output, and  $y_i$  is the actual output.

**Regularized Loss with L2 Regularization:**

$$\mathcal{L} = \mathcal{L}_{MSE} + \lambda \sum_l ||W_l||^2$$

where  $\lambda$  is the regularization strength, and  $||W_l||^2$  is the squared sum of the weights.

During backpropagation, the weight updates follow:

$$W_l \leftarrow W_l - \eta \left( \frac{\partial \mathcal{L}_{MSE}}{\partial W_l} + 2\lambda W_l \right)$$

where  $\eta$  is the learning rate. The additional term  $2\lambda W_l$  pulls weights toward zero, preventing overfitting.

**Impact of L2 Regularization on Weight Updates:**

- The term  $-2\lambda W_l$  acts like a friction force, pushing the weights toward zero.
- This discourages large weight values, making the model simpler and more generalizable to unseen data.
- It does **not** force weights to exactly zero (unlike L1 regularization, which induces sparsity), but it reduces their magnitude gradually.
- The impact depends on  $\lambda$ :
  - If  $\lambda$  is too small, regularization is weak, and overfitting may still occur.
  - If  $\lambda$  is too large, regularization is too strong, and the model may **underfit** (failing to learn meaningful patterns).

Task 3: Draw a computational graph (examples: 1, 2, 3) for the autoencoder. Automated computational graph generators are not allowed.

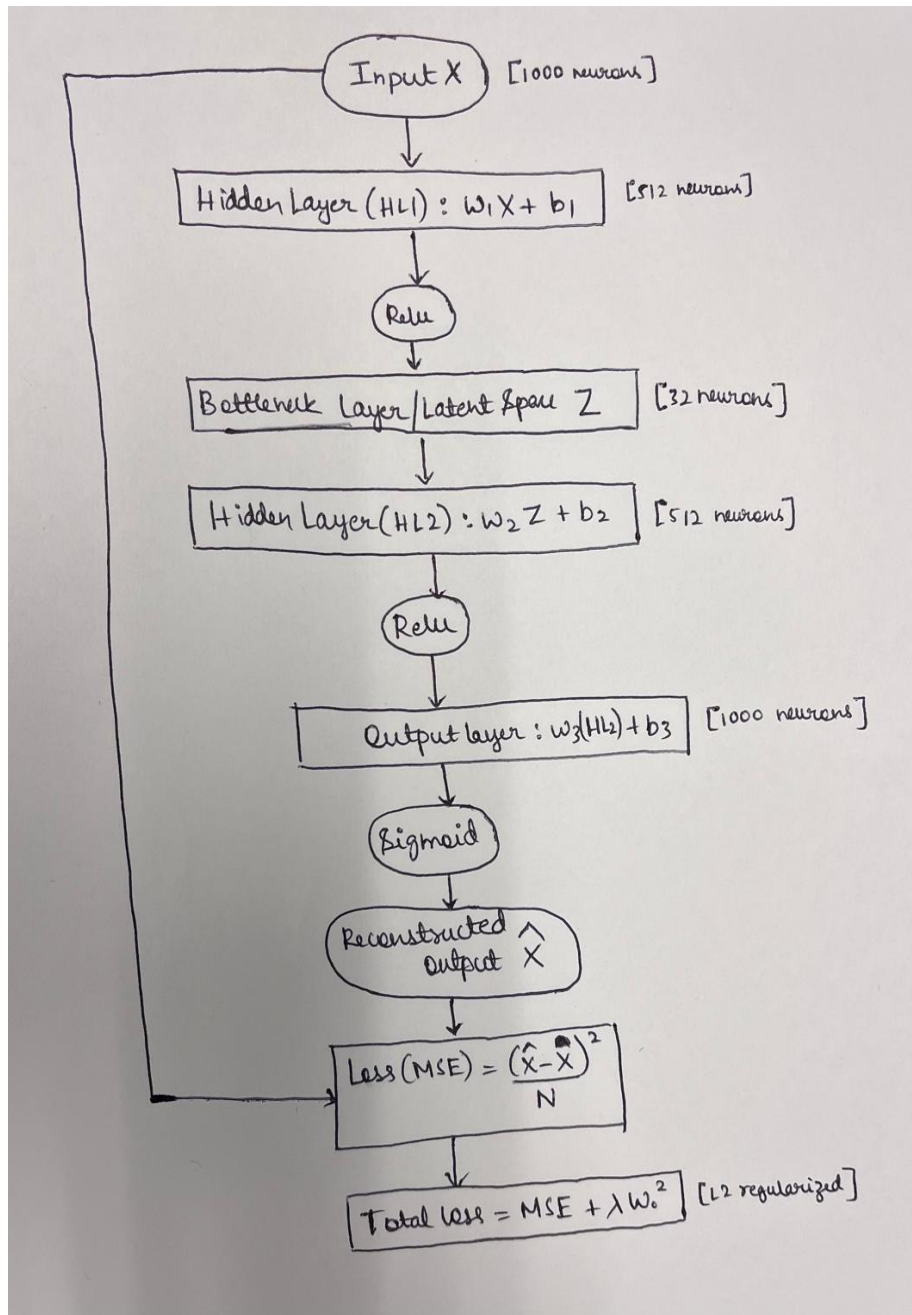


Figure 1: Computational graph of the autoencoder

**Task 4: Discuss potential challenges and limitations (at least 4) of using autoencoders for anomaly detection in manufacturing. Consider data, anomaly definition, model assumptions, and deployment.**

1. **Defining Anomalies:** Anomalies are often rare, making it difficult to train the model effectively. The model may struggle to generalize if the definition of anomalies is unclear or changes over time.
2. **Data Quality and Preprocessing:** The autoencoder assumes that normal data patterns are well-represented in training. Poor data quality (sensor noise, missing values) can lead to inaccurate reconstructions and false positives.
3. **Overfitting to Normal Data:** The model may learn to reconstruct normal data too well, leading to low reconstruction errors even for slightly anomalous data.
4. **Deployment and Real-time Performance:** Real-time anomaly detection requires low latency and integration with existing systems which can be challenging.

**Task 5: Propose potential improvements or extensions (at least 4) to the system to enhance its effectiveness in detecting anomalies. Consider architecture, loss, preprocessing, integration.**

1. **Architectural Enhancements:** Using Variational Autoencoders (VAEs) or Convolutional Autoencoders to improve feature extraction.
2. **Improved Loss Function:** Instead of plain MSE, use contrastive loss or weighted loss functions that emphasize anomalies, making the model more sensitive to rare defects.
3. **Preprocessing and Feature Engineering:** Apply dimensionality reduction (PCA, t-SNE) and denoising techniques to remove sensor noise before feeding data into the autoencoder.
4. **Integration with Other Methods:** Combine autoencoders with statistical methods (Z-score, IQR) or machine learning classifiers (SVM, Isolation Forest) to enhance anomaly detection accuracy.

## Part I. || Transformers & Self-Attention

### Task 1:

#### 1. Mathematical Operations in Self-Attention :

Given an input sequence,

$$x = (x_1, x_2, \dots, x_n);$$

the self-attention mechanism proceeds in the following stages :

##### (a) Linear Transformations

1. Query ( $q$ ) : Queries represent the vector that "asks" about relationships with other elements. Each input  $x_i$  vector is mapped to a query vector  $q_i$ .

$$q_i = W_q x_i + b_q,$$

where  $W_q$  is the query weight matrix and  $b_q$  is the corresponding bias vector.

2. Key ( $k$ ) : Keys represent the vector that "answers" the queries, signalling relevance to each element. Each  $x_i$  is also mapped to a key vector  $k_i$ .

$$k_i = W_k x_i + b_k,$$

where  $W_k$  is the key weight matrix and  $b_k$  is its bias.

3. Value ( $v$ ) : Values carry the actual content information of each element that will be weighted by attention scores.

$$v_i = W_v x_i + b_v,$$

Each  $W_v$  is the value weight matrix and  $b_v$  is its bias.



Typically,  $W_q$ ,  $W_k$  and  $W_v$  are all (input dim)  $\times$   $d$  matrices (where  $d$  might be the embedding dimension or a chosen "attention dimension"), and the biases are  $d$ -dimensional vectors. These learnable parameters let the model transform each input query, key & value subspaces tailored for attention.

## (b) Scaled Dot-Product Attention

Once the queries  $q_i$  and keys  $k_j$  are computed, the attention score between element  $i$  and element  $j$  is obtained via the dot product, scaled by  $\sqrt{d_k}$  (the square root of  $k$  dim):

$$\text{score}(i, j) = \frac{q_i \cdot k_j}{\sqrt{d_k}}$$

- without scaling, the dot product  $q_i \cdot k_j$  could grow large in magnitude as  $d_k$  increases, leading to large gradients & making the softmax distribution overly peaked. Dividing by  $\sqrt{d_k}$  normalizes these scores, stabilizing training & helping the model distribute attention more effectively.
- Normalization: after computing the raw scores  $\text{score}(i, j)$ , we typically apply a softmax across  $j$  (for each fixed  $i$ ) to obtain attention weights:

$$a_{ij} = \text{softmax}_j(\text{score}(i, j)) = \frac{\exp(\text{score}(i, j))}{\sum_j \exp(\text{score}(i, j))}$$

## (c) Weighted Sum:

Using the attention weights  $a_{ij}$ , we compute a weighted sum of the value vectors  $v_j$ . For each position  $i$ :

$$z_i = \sum_{j=1}^N a_{ij} v_j$$

Intuitively, the model "looks at" all the values in the sequence (through  $x_{ij}$ ), but places greater weight on the most relevant elements (i.e., those with higher attention scores).

#### (d) Optional Output Transformation:-

Finally, an optional linear transformation can be applied to each  $z_i$ . Commonly, we write:

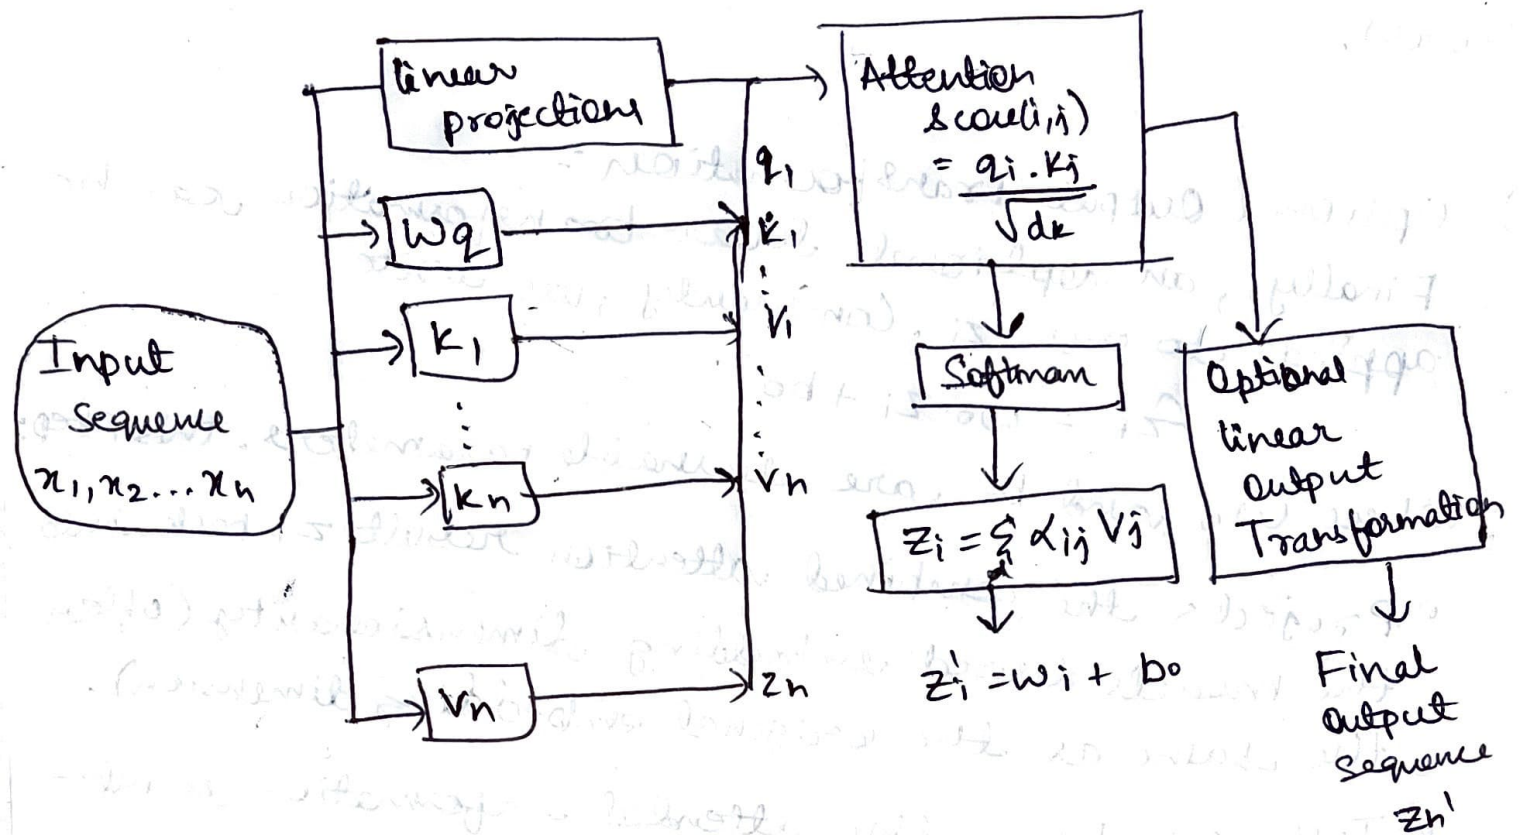
$$\hat{z}_i = w_o z_i + b_o,$$

where  $w_o$  and  $b_o$  are learnable parameters. This step:

- Projects the combined attention result  $z_i$  back into the model's desired embedding dimensionality (often the same as the original embedding dimension).
- Influences how the attended information is integrated into the broader model's representation. If this projection is omitted,  $z_i$  itself can be used directly, but typically the model includes this learnable output matrix to enhance expressive power.
- Including this step leads to better overall performance and flexibility within Transformer architectures.



## 2. Computational Graph :





### **Contribution Table:**

| Team Member               | Assignment Part    | Contribution (%) |
|---------------------------|--------------------|------------------|
| Ruthvik Vasantha Kumar    | 1,2,3,4, all bonus | 50%              |
| Shreyas Bellary Manjunath | 1,2,3,4, all bonus | 50%              |